

Designing Hardware to Improve Your Software

A Use Case in Sound



About Me

- Undergrad in EECS @ MIT
- Worked at JetBrains, NVIDIA, and Jump Trading
- Passionate about Linux, audio, and embedded systems



Goals for the talk




1. How do operating systems work?
2. How does sound work on Linux?
3. How can we build our own hardware with FPGAs?
4. Describe the journey I took of building something complex from scratch



Computation Resources



Hardware Peripherals







Scheduler



time





```
POST / HTTP/1.1
Host: developer.mozilla.org
User-Agent: curl/8.6.0
Accept: */*
Content-Type: application/json
Content-Length: 345
```

```
{
  "data": "ABC123"
}
```



```
HTTP/1.1 403 Forbidden
Server: Apache
Date: Fri, 21 Jun 2024 12:52:39 GMT
Content-Length: 678
Content-Type: text/html
Cache-Control: no-store
```

```
<!DOCTYPE html>
<html lang="en">
(more data...)
```


Register	Argument User Space	Argument Kernel Space
%rax	Not Used	System Call Number
%rdi	Arguement 1	Arguement 1
%rsi	Arguement 2	Arguement 2
%rdx	Arguement 3	Arguement 3
%r10	Not Used	Arguement 4
%r8	Arguement 5	Arguement 5
%r9	Arguement 6	Arguement 6
%rcx	Arguement 4	Destroyed
%r11	Not Used	Destroyed

A L S A

Advanced Linux Sound Architecture

- Created by Jaroslav Kysela in 1998
- Merged into Linux 2.5 in 2002
- Framework for writing sound card drivers
- Userspace API for interacting with sound cards

ALSA
Apps

JACK

Pulse
Audio

Pipe
Wire

ALSA Userspace C Library

ALSA Syscall Interface

ALSA Sound Card Driver

Userspace

Kernel



Card 0

/dev/snd/controlC0

Device 0

/dev/snd/pcmC0D0p

Headphones

Device 1

/dev/snd/pcmC0D1c

Microphone

Device 2

/dev/snd/pcmC0D2p

HDMI Out

What do you mean “my sound card is just a bunch of files”?


Well, not just your sound card. Everything is.



`/dev/input/by-id/<id>`




`/dev/sda`




`/dev/mem`

```
open(path, flags)
```



“File Descriptor”



```
close(fd)
```


`ioctl` (`fd`, `cmd`, [`data`])

input/output control

file descriptor

command

optional data

The diagram illustrates the `ioctl` function signature. The function name `ioctl` is in light gray and has a white bracket underneath it labeled "input/output control". The opening parenthesis is also light gray. The first argument `fd` is in yellow and has a yellow bracket above it labeled "file descriptor". The second argument `cmd` is in red and has a red bracket below it labeled "command". The third argument is an optional array `[data]`, where the opening square bracket is light gray, `data` is in green, and the closing square bracket is light gray. A green bracket above the `data` is labeled "optional data". The closing parenthesis is light gray.

(Excerpt from /usr/include/sound/asound.h)

```
#define SNDRV_PCM_IOCTL_WRITE16FS _IOW('A', 0x52, struct snd_xfern)
```

01

000000000011000

010000001

01010010

1075331410

(in decimal)

◆ snd_pcm_open()

```
int snd_pcm_open ( snd_pcm_t **      pcmp,  
                  const char *        name,  
                  snd_pcm_stream_t stream,  
                  int                 mode  
                  )
```

Opens a PCM.

Parameters

pcmp Returned PCM handle
name ASCII identifier of the PCM handle
stream Wanted stream
mode Open mode (see [SND_PCM_NONBLOCK](#), [SND_PCM_ASYNC](#))

```
// Open card  
ioctl("/dev/snd/controlC0", SNDRV_CTL_IOCTL_CARD_INFO)  
ioctl("/dev/snd/controlC0", SNDRV_CTL_IOCTL_PVERSION)  
ioctl("/dev/snd/controlC0", SNDRV_CTL_IOCTL_PCM_PREFER_SUBDEVICE)  
  
// Open PCM  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_INFO)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_PVERSION)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_USER_PVERSION)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_TTSTAMP)
```

◆ snd_pcm_hw_params_any()

```
int snd_pcm_hw_params_any ( snd_pcm_t * pcm,  
                           snd_pcm_hw_params_t * params  
                           )
```

Fill params with a full configuration space for a PCM.

Parameters

pcm PCM handle

params Configuration space

The configuration space will be filled with all possible ranges for the PCM device.

Note that the configuration space may be constrained by the currently installed configuration on the PCM device. To remove any constraints, free the configuration with `snd_pcm_hw_free` first.

```
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_HW_REFINE)
```

◆ snd_pcm_hw_params()

```
int snd_pcm_hw_params ( snd_pcm_t *      pcm,  
                       snd_pcm_hw_params_t * params  
                       )
```

Install one PCM hardware configuration chosen from a configuration space and `snd_pcm_prepare` it.

Parameters

pcm PCM handle
params Configuration space definition container

Returns

0 on success otherwise a negative error code

The configuration is chosen fixing single parameters in this order: first access, first format, first subformat, min channels, min rate, min period time, max buffer size, min tick time. If no mutually compatible set of parameters can be chosen, a negative error code will be returned.

After this call, `snd_pcm_prepare()` is called automatically and the stream is brought to `SND_PCM_STATE_PREPARED` state.

The hardware parameters cannot be changed when the stream is running (active). The software parameters can be changed at any time.

The configuration space will be updated to reflect the chosen parameters.

```
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_HW_REFINE)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_HW_PARAMS)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_PREPARE)
```

◆ snd_pcm_sw_params()

```
int snd_pcm_sw_params ( snd_pcm_t * pcm,  
                        snd_pcm_sw_params_t * params  
                        )
```

Install PCM software configuration defined by params.

Parameters

pcm PCM handle

params Configuration container

Returns

0 on success otherwise a negative error code

The software parameters can be changed at any time. The hardware parameters cannot be changed when the stream is running (active).

The function is thread-safe when built with the proper option.

```
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_SW_PARAMS)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_PREPARE)
```

♦ snd_pcm_writei()

```
snd_pcm_sframes_t snd_pcm_writei ( snd_pcm_t *      pcm,  
                                   const void *      buffer,  
                                   snd_pcm_uframes_t size  
                                   )
```

Write interleaved frames to a PCM.

Parameters

pcm PCM handle
buffer frames containing buffer
size frames to be written

Returns

a positive number of frames actually written otherwise a negative error code

```
// Start filling up audio buffer  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_WRITEN_FRAMES)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_HWSYNC)  
  
// Once enough data is sent, begin playback  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_START)  
  
// Continue sending audio data as playback continues  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_WRITEN_FRAMES)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_HWSYNC)
```

◆ snd_pcm_drain()

```
int snd_pcm_drain ( snd_pcm_t * pcm )
```

Stop a PCM preserving pending frames.

Parameters

pcm PCM handle

Returns

0 on success otherwise a negative error code

```
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_SW_PARAMS)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_DRAIN)
```


◆ snd_pcm_close()

```
int snd_pcm_close ( snd_pcm_t * pcm )
```

close PCM handle

Parameters

pcm PCM handle

Returns


0 on success otherwise a negative error code

Closes the specified PCM handle and frees all associated resources.

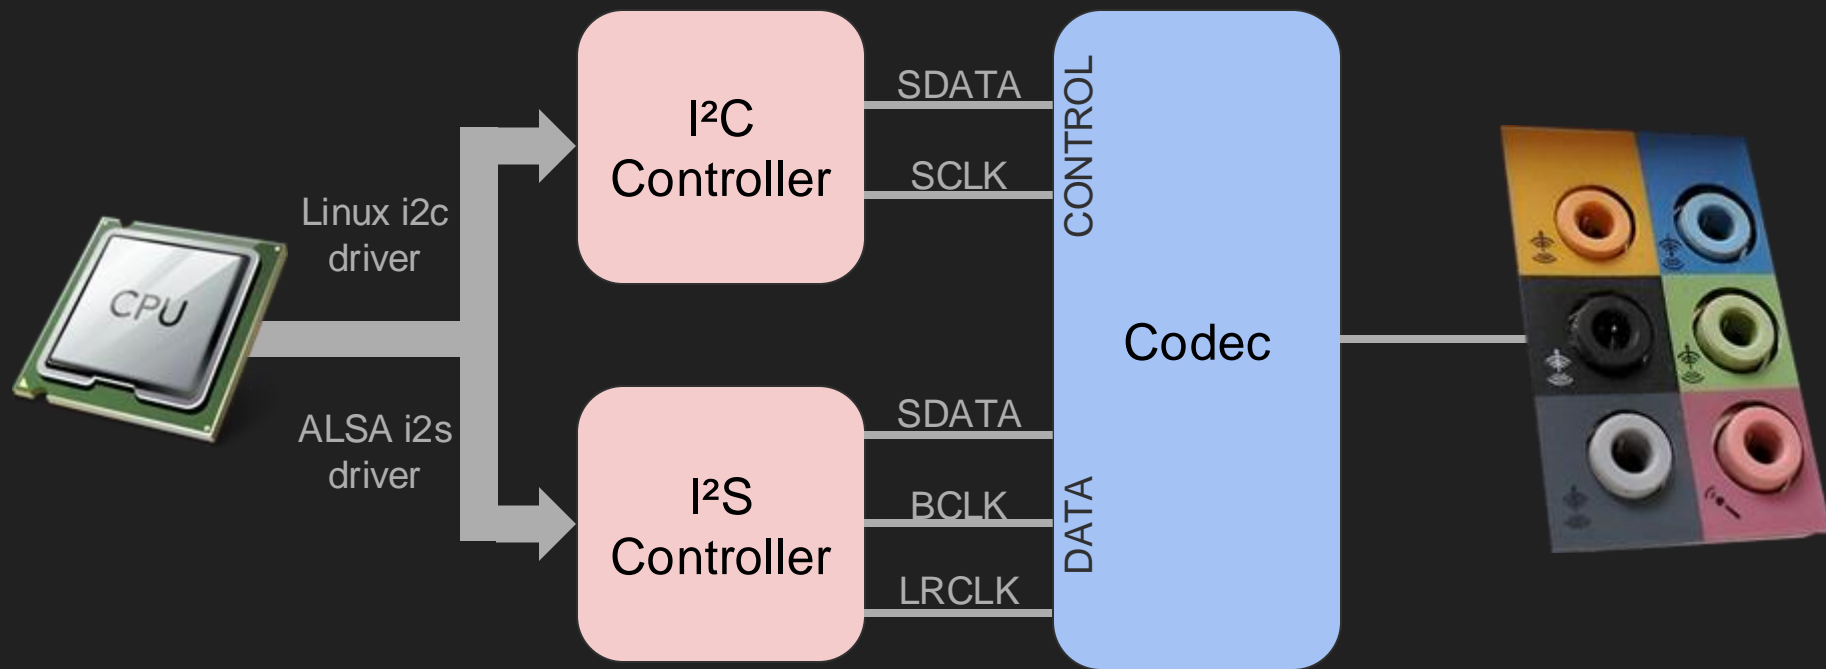
```
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_DROP)  
ioctl("/dev/snd/pcmC0D0p", SNDRV_PCM_IOCTL_HW_FREE)
```


How many ioctl's does ALSA define?

91



```
struct snd_pcm_ops {
    int (*open)(struct snd_pcm_substream *substream);
    int (*close)(struct snd_pcm_substream *substream);
    int (*ioctl)(struct snd_pcm_substream * substream,
        unsigned int cmd, void *arg);
    int (*hw_params)(struct snd_pcm_substream *substream,
        struct snd_pcm_hw_params *params);
    int (*hw_free)(struct snd_pcm_substream *substream);
    int (*prepare)(struct snd_pcm_substream *substream);
    int (*trigger)(struct snd_pcm_substream *substream, int cmd);
    int (*sync_stop)(struct snd_pcm_substream *substream);
    snd_pcm_uframes_t (*pointer)(struct snd_pcm_substream *substream);
    int (*get_time_info)(struct snd_pcm_substream *substream,
        struct timespec64 *system_ts, struct timespec64 *audio_ts,
        struct snd_pcm_audio_tstamp_config *audio_tstamp_config,
        struct snd_pcm_audio_tstamp_report *audio_tstamp_report);
    int (*fill_silence)(struct snd_pcm_substream *substream, int channel,
        unsigned long pos, unsigned long bytes);
    int (*copy)(struct snd_pcm_substream *substream, int channel,
        unsigned long pos, struct iov_iter *iter, unsigned long bytes);
    struct page *(*page)(struct snd_pcm_substream *substream,
        unsigned long offset);
    int (*mmap)(struct snd_pcm_substream *substream, struct vm_area_struct *vma);
    int (*ack)(struct snd_pcm_substream *substream);
};
```





What if we want more?

F P G A ield rogrammable ate rray


```
library ieee;
use ieee.std_logic_1164.all;

entity blink_led is
port (
    i_clk : in std_logic;
    o_led : out std_logic;
);
end blink_led;

architecture behavioral of blink_led is
    constant CLKS_PER_SEC : natural := 100000000;
    signal counter : natural := 0;
    signal led : std_logic := '0';
begin
    blink_proc : process(i_clk)
    begin
        if rising_edge(i_clk) then
            if counter < CLKS_PER_SEC then
                counter <= counter + 1;
            else
                led <= not led;
                counter <= 0;
            end if;
        end if;
    end process;

    o_led <= led;
end behavioral;
```

Hardware Description Language



Programming Language

- Program begins at a specific point
- Execution proceeds step-by-step
- Data is shared by putting it into CPU registers or RAM
- Program exits once it finishes its work

Inherently **serial**

Hardware Description Language

- Entire design starts simultaneously
- All logic executes simultaneously
- Data is shared by running “wires” to every place where the data needs to be read
- Design runs forever

Inherently **parallel**

High Level
Synthesis
(HLS)

Bluespec

Chisel

VHDL

Verilog

Register Transfer Level

VHDL Crash Course

Entity

Architecture

i_clk

i_input1

i_input2

signal1

signal2

o_output1

o_output2

o_output3

```
library ieee;
    use ieee.std_logic_1164.all;

entity my_entity is
    port (
        i_clk      : in  std_logic;
        i_input1   : in  std_logic;
        i_input2   : in  std_logic;
        o_output1  : out std_logic;
        o_output2  : out std_logic;
        o_output3  : out std_logic;
    );
end my_entity;

architecture my_architecture of my_entity is
    signal signal1 : std_logic := '0';
    signal signal2 : std_logic := '0';
begin
    -- Implementation goes here
end my_architecture;
```

Dataflow Modelling

- Combinatorial logic drives outputs
- Uses concurrent signal assignment statements

```
library ieee;
    use ieee.std_logic_1164.all;

entity full_adder is
    port (
        i_a      : in  std_logic;
        i_b      : in  std_logic;
        i_carry   : in  std_logic;
        o_sum     : out std_logic;
        o_carry   : out std_logic;
    );
end full_adder;

architecture dataflow of full_adder is
begin
    o_sum  <= i_a xor i_b xor i_carry;

    o_carry <= (i_a and i_b)      or
               (i_a and i_carry) or
               (i_b and i_carry);
end dataflow;
```

```
library ieee;
    use ieee.std_logic_1164.all;

-- Assume this is defined in another package:
-- type State_t is (INIT, WAIT_FOR_DATA, PROCESSING, DONE);

entity state_display is
    port (
        i_state : in  State_t;
        o_leds  : out std_logic_vector(3 downto 0);
    );
end state_display;

architecture dataflow of state_display is
begin
    with state select
        o_leds <= "1000" when STATE = INIT,
                  "0100" when STATE = WAIT_FOR_DATA,
                  "0010" when STATE = PROCESSING,
                  "0001" when STATE = DONE,
                  "0000" when others;
end dataflow;
```

Behavioral Modelling

- Event-driven state machine drives outputs
- Uses **process** statements

```
library ieee;
use ieee.std_logic_1164.all;

entity blink_led is
    port (
        i_clk : in  std_logic;
        o_led : out std_logic;
    );
end blink_led;

architecture behavioral of blink_led is
    constant CLKS_PER_SEC : natural := 100000000;
    signal counter : natural := 0;
    signal led : std_logic := '0';
begin

    blink_proc : process(i_clk)
    begin
        if rising_edge(i_clk) then
            if counter < CLKS_PER_SEC then
                counter <= counter + 1;
            else
                led <= not led;
                counter <= 0;
            end if;
        end if;
    end process;

    o_led <= led;

end behavioral;
```

Structural Modelling

- Instantiated sub-entities drive outputs
- Uses **component instantiation** statements

```
library ieee;
use ieee.std_logic_1164.all;


-- Bring in component definitions from our package
library work;
use work.all;

entity two_bit_adder is
port (
    i_a      : in  std_logic_vector(0 to 1);
    i_b      : in  std_logic_vector(0 to 1);
    o_sum     : out std_logic_vector(0 to 1);
    o_carry   : out std_logic;
);
end two_bit_adder;


architecture structural of top is
    signal carry : std_logic := '0';
begin
    first_full_adder : work.full_adder
        port map (
            i_a => i_a(0),
            i_b => i_b(0),
            i_carry => '0',
            o_sum  => o_sum(0),
            o_carry => carry
        );

    second_full_adder : work.full_adder
        port map (
            i_a      => i_a(1),
            i_b      => i_b(1),
            i_carry  => carry,
            o_sum    => o_sum(1),
            o_carry  => o_carry
        );
end structural;
```


Vision




Vision



Phase 1 Gameplan

- Buy a book on VHDL
- Download vivado & build command-line workflow
 - Requires writing TCL scripts for IP generation, synthesis, and programming
- Write `spdif_tx` and `spdif_rx` entities in VHDL
- Create circuits to shift S/PDIF coax from TTL <-> 1V



200 ns/div

100 MHz

-800 mV

Aut

of 64

Instruments

Auto setup

Open

Save

Print

Full









#12

Joined Nov 30, 2010

18,224

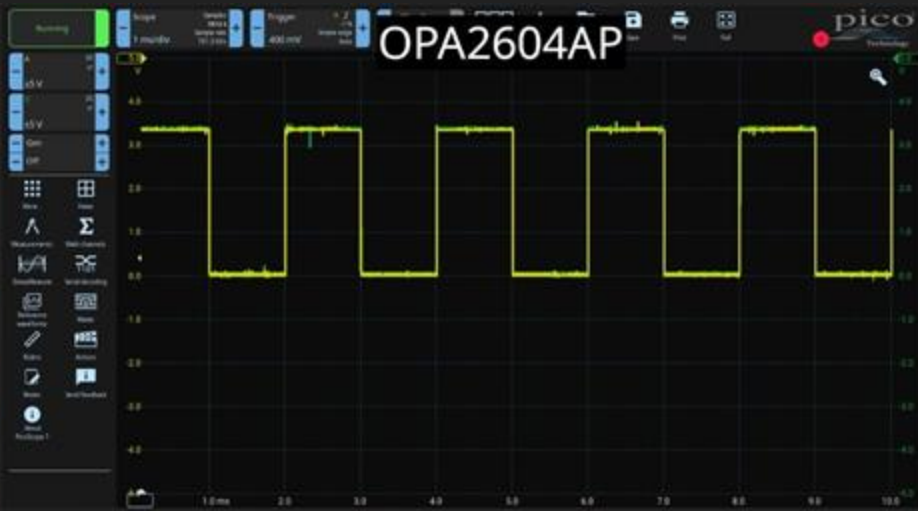
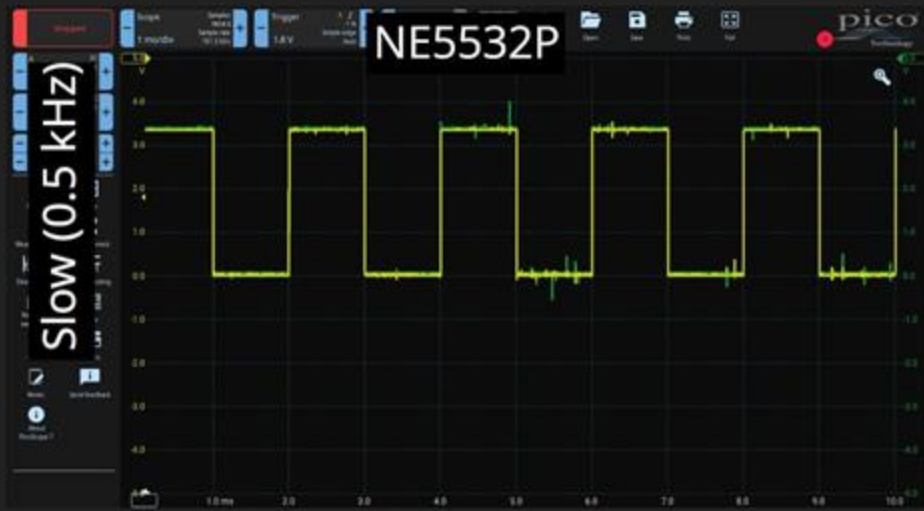
Apr 22, 2016

#2

One way to achieve fractional gain is to use 2 resistors and no amplifier. Another way is to use the inverting configuration and make the feedback resistor smaller than the input resistor. If you don't like the polarity inversion, add another inverting gain stage with a gain of one.







spdif_tx



```
graph LR; A[spdif_tx] --> B[Observed Result: No audio transmitted];
```

Observed Result:





No audio
transmitted

Phase 2 Gameplan

- Buy a book on Linux device drivers
- Find an example ALSA driver online
 - `dummy.c`, ~1200 lines of C code without even a single comment
- Refactor into multiple source files
- Study the code until I understood it

Phase 3 Gameplan

- Build ability to send/recv ethernet from FPGA



MII

GMII

RMII

RGMII

Phase 3 Gameplan

- Build ability to send/recv ethernet from FPGA
- Build ability to send/recv ethernet from driver
- Design wire protocol
- Build out device handshake & heartbeating
- Copy audio data into ethernet frames
- Build audio buffer in FPGA to store samples
- Connect audio buffer to spdif_rx and spdif_tx

Demo

